

Activity 1

Checking the MAC address of the network card

```
ifconfig

# -c flag colour the output,
# a flag means printing of all interfaces
ip -c a

# we can check the MAC address of a specific interface using the NetworkManager
tool
nmcli dev show ens160

MAC address for ens160 interface is: 00:0c:29:7d:1f:f2
```

```
rocky@rocky-server-3:~$ ifconfig
ens160: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.1.42 netmask 255.255.255.0 broadcast 192.168.1.255
    inet6 fe80::20c:29ff:fe7d:1ff2 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:7d:1f:f2 txqueuelen 1000 (Ethernet)
    RX packets 18859 bytes 22647557 (21.5 MiB)
    RX errors 0 dropped 892 overruns 0 frame 0
    TX packets 8363 bytes 974739 (951.8 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[rocky@rocky-server-3 ~]$ ip -c a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens160: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 00:0c:29:7d:1f:f2 brd ff:ff:ff:ff:ff:ff
    altname enp3s0
    inet 192.168.1.42/24 brd 192.168.1.255 scope global dynamic noprefixroute ens160
        valid_lft 81794sec preferred_lft 81794sec
    inet6 fe80::20c:29ff:fe7d:1ff2/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
[rocky@rocky-server-3 ~]$ nmcli dev show ens160
GENERAL.DEVICE: ens160
GENERAL.TYPE: ethernet
GENERAL.HWADDR: 00:0C:29:7D:1F:F2
GENERAL.MTU: 1500
GENERAL.STATE: 100 (connected)
GENERAL.CONNECTION: ens160
GENERAL.CON-PATH: /org/freedesktop/NetworkManager/ActiveConnection/2
WIRED-PROPERTIES.CARRIER: on
IP4.ADDRESS[1]: 192.168.1.42/24
IP4.GATEWAY: 192.168.1.1
IP4.ROUTE[1]: dst = 192.168.1.0/24, nh = 0.0.0.0, mt = 100
IP4.ROUTE[2]: dst = 0.0.0.0/0, nh = 192.168.1.1, mt = 100
IP4.DNS[1]: 192.168.1.1
IP6.ADDRESS[1]: fe80::20c:29ff:fe7d:1ff2/64
IP6.GATEWAY: --
IP6.ROUTE[1]: dst = fe80::/64, nh = ::, mt = 1024
[rocky@rocky-server-3 ~]$ _
```

Checking local ARP table

```
arp -ni ens160 # i flag allows us to select a specific interface, n flag show
numerical addresses instead of symbolic host names
ip neigh
```

```

[rocky@rocky-server-3 ~]$ arp -ni ens160
Address          HWtype  HWaddress      Flags Mask    Iface
192.168.1.1      ether   40:b0:76:c3:97:b8  C           ens160
192.168.1.19     ether   70:85:c2:68:14:e8  C           ens160
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ ip neigh
192.168.1.1 dev ens160 lladdr 40:b0:76:c3:97:b8 STALE
192.168.1.19 dev ens160 lladdr 70:85:c2:68:14:e8 REACHABLE
[rocky@rocky-server-3 ~]$ _

```

Checking network interfaces and associated IP addresses

```

ifconfig
ip a
nmcli

```

```

[rocky@rocky-server-3 ~]$ ifconfig
ens160: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.1.42 netmask 255.255.255.0 broadcast 192.168.1.255
    inet6 fe80::20c:29ff:fe7d:1ff2 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:7d:1f:f2 txqueuelen 1000 (Ethernet)
    RX packets 23979 bytes 23280218 (22.2 MiB)
    RX errors 0 dropped 1751 overruns 0 frame 0
    TX packets 12171 bytes 1757529 (1.6 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 52 bytes 3328 (3.2 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 52 bytes 3328 (3.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[rocky@rocky-server-3 ~]$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens160: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 00:0c:29:7d:1f:f2 brd ff:ff:ff:ff:ff:ff
    altname enp3s0
    inet 192.168.1.42/24 brd 192.168.1.255 scope global dynamic noprefixroute ens160
        valid_lft 77421sec preferred_lft 77421sec
    inet6 fe80::20c:29ff:fe7d:1ff2/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
[rocky@rocky-server-3 ~]$ _

```

```

[rocky@rocky-server-3 ~]$ nmcli
ens160: connected to ens160
    "VMware VMXNET3"
    ethernet (vmxnet3), 00:0C:29:7D:1F:F2, hw, mtu 1500
    ip4 default
    inet4 192.168.1.42/24
    route4 192.168.1.0/24 metric 100
    route4 default via 192.168.1.1 metric 100
    inet6 fe80::20c:29ff:fe7d:1ff2/64
    route6 fe80::/64 metric 1024

lo: connected (externally) to lo
    "lo"
    loopback (unknown), 00:00:00:00:00:00, sw, mtu 65536
    inet4 127.0.0.1/8
    inet6 ::1/128
    route6 ::1/128 metric 256

DNS configuration:
    servers: 192.168.1.1
    interface: ens160

Use "nmcli device show" to get complete information about known devices and
"nmcli connection show" to get an overview on active connection profiles.

Consult nmcli(1) and nmcli-examples(7) manual pages for complete usage details.
[rocky@rocky-server-3 ~]$ _

```

Checking the routing table

```
ip r # r flag shows table routes
route -n # n flag show numerical addresses instead of symbolic host names
netstat -rn # n flag show numerical addresses instead of symbolic host names, r
flag shows routing tables
```

```

[rocky@rocky-server-3 ~]$ ip r
default via 192.168.1.1 dev ens160 proto dhcp src 192.168.1.42 metric 100
192.168.1.0/24 dev ens160 proto kernel scope link src 192.168.1.42 metric 100
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ route -n
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0        192.168.1.1    0.0.0.0         UG    100    0      0 ens160
192.168.1.0    0.0.0.0        255.255.255.0   U     100    0      0 ens160
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ netstat -rn
Kernel IP routing table
Destination    Gateway         Genmask         Flags MSS Window  irtt Iface
0.0.0.0        192.168.1.1    0.0.0.0         UG      0      0      0 ens160
192.168.1.0    0.0.0.0        255.255.255.0   U       0      0      0 ens160
[rocky@rocky-server-3 ~]$ _

```

Checking list of open (listening) TCP ports

```

sudo ss -tlp # t flag means TCP, l flag means listening ports, p flag means port
sudo lsof -iTCP -sTCP:LISTEN # List all TCP connections with state LISTEN
sudo nmap -sT 192.168.1.42 # -sT means TCP connect scan
sudo netstat -tlnp # t means TCP, l listening, p port, n show numerical addresses
instead of symbolic host names

```

```

[rocky@rocky-server-3 ~]$ sudo ss -tlp
State      Recv-Q    Send-Q     Local Address:Port      Peer Address:Port      Process
LISTEN     0          128             0.0.0.0:ssh              0.0.0.0:*                users:((("sshd",pid=822,fd=3))
LISTEN     0          128             :::ssh                   :::*                     users:((("sshd",pid=822,fd=4))
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ sudo lsof -iTCP -sTCP:LISTEN
COMMAND PID USER   FD   TYPE DEVICE SIZE/OFF NODE NAME
sshd    822 root    3u    IPv4  22486   0t0  TCP *:ssh (LISTEN)
sshd    822 root    4u    IPv6  22488   0t0  TCP *:ssh (LISTEN)
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ sudo nmap -sT 192.168.1.42
Starting Nmap 7.91 ( https://nmap.org ) at 2023-11-17 22:08 CET
Nmap scan report for rocky-server-3 (192.168.1.42)
Host is up (0.00016s latency).
Not shown: 999 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh
Nmap done: 1 IP address (1 host up) scanned in 0.06 seconds
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ sudo netstat -tlnp
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address   Foreign Address  State       PID/Program name
tcp        0      0 0.0.0.0:22      0.0.0.0:*        LISTEN      822/sshd: /usr/sbin
tcp6       0      0 :::22           :::*             LISTEN      822/sshd: /usr/sbin
[rocky@rocky-server-3 ~]$ _

```

Activity 2

Subtask 1

```

# c flag sending count ECHO_REQUEST packets.
ping 8.8.8.8 -c 5

# Output
# 0% packets loss
# RTT between my PC and google server (8.8.8.8) = min: 14.829, max: 15.401, avg:
15.170

```

```

[rocky@rocky-server-3 ~]$ ping 8.8.8.8 -c 5
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=119 time=14.9 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=119 time=15.3 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=119 time=15.4 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=119 time=15.3 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=119 time=14.8 ms

--- 8.8.8.8 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4007ms
rtt min/avg/max/mdev = 14.829/15.170/15.401/0.235 ms
[rocky@rocky-server-3 ~]$ _

```

Subtask 2

`tracert -n 8.8.8.8` # n flag means that tracert doesn't try to map IP addresses to host names.

`mtr -rn -c5 8.8.8.8`

r flag puts mtr into report mode

n flag display numeric IP numbers and not try to resolve the host names

c flag set the number of pings sent to determine both the machines on the network and the reliability of those machines.

```
rocky@rocky-server-3:~$ tracert -n 8.8.8.8
tracert to 8.8.8.8 (8.8.8.8), 30 hops max, 60 byte packets
 1 192.168.1.1  1.676 ms  1.402 ms  1.304 ms
 2 172.16.0.2   3.474 ms  3.403 ms  3.321 ms
 3 10.64.0.2    40.543 ms 40.459 ms 40.377 ms
 4 10.64.0.1    2.983 ms  3.135 ms  3.058 ms
 5 185.48.10.148 7.065 ms  6.994 ms  7.092 ms
 6 192.168.51.1 12.258 ms 12.100 ms 11.810 ms
 7 195.182.219.69 12.492 ms 10.543 ms 10.704 ms
 8 142.250.37.209 15.224 ms 108.170.250.209 15.386 ms 108.170.250.193 16.064 ms
 9 172.253.68.31 11.844 ms 216.239.40.43 11.332 ms 142.250.224.91 11.187 ms
10 8.8.8.8      10.448 ms 10.823 ms 10.686 ms
[rocky@rocky-server-3 ~]$
[rocky@rocky-server-3 ~]$ mtr -rn -c5 8.8.8.8
Start: 2023-11-18T12:47:43+0100
HOST: rocky-server-3
  Loss%  Snt  Last  Avg  Best  Wrst  StDev
 1 |-- 192.168.1.1      0.0%    5   1.4   1.3   0.9   1.5   0.2
 2 |-- 172.16.0.2       0.0%    5   3.0   3.0   2.8   3.1   0.1
 3 |-- 10.64.0.2        0.0%    5   3.1   3.2   2.6   3.7   0.4
 4 |-- 10.64.0.1        0.0%    5   3.5   3.7   3.5   4.0   0.2
 5 |-- 185.48.10.148    0.0%    5   7.3   7.2   7.1   7.3   0.1
 6 |-- 192.168.51.1     0.0%    5  12.1  12.1  12.0  12.3   0.1
 7 |-- 195.182.219.69   0.0%    5  12.2  12.1  12.1  12.2   0.1
 8 |-- 108.170.250.193  0.0%    5  16.3  15.7  14.6  16.7   0.8
 9 |-- 108.170.234.101  0.0%    5  12.1  12.1  11.8  12.3   0.2
10 |-- 8.8.8.8          0.0%    5  15.5  15.6  15.4  15.8   0.2
[rocky@rocky-server-3 ~]$ _
```

Activity 3

Subtask 1

192.168.0.0/26

Netmask 255.255.255.192

Network address 192.168.0.0

Broadcast address 192.168.0.63

Number of hosts in the subnet: 62

- 1. Netmask calculation:**
CIDR 26
192.168.0.0/26
Each octet is 8 bits, CIDR is 26, it means that we need to calculate only the last octet
11111111. 11111111. 11111111.?
In 4th octet we left with two bits
11111111. 11111111. 11111111.11000000
Power of two: 128 64 32 16 8 4 2 1
128 + 64 = 192
NM 255.255.255.192
- 2. Network address calculation:**
192.168.0.0/26
The network address in the network is the first one, so we fill it with zeros or we can follow the rules below:
0 AND 0 = 0
1 AND 0 = 0
0 AND 1 = 0
1 AND 1 = 1
IP 192.168.0.0 = last octet is 0 0 0 0 0 0 0 0
NA 192.168.0.0
- 3. Broadcast address calculation:**
192.168.0.0/26
The broadcast address in the network is the last one, so we fill it with ones
BA 192.168.0.0 0 1 1 1 1 1 1 = 63
BA 192.168.0.63
- 4. Calculating the number of hosts in a subnet:**
192.168.0.0/26
IP address is 32-bit number. CIDR is 26, it means:
 $32 - 26 = 6$
 $2^6 = 64$
We have a total of 64 addresses, but we also need to subtract the network address and the broadcast address.
 $64 - 2 = 62$
Number of hosts in the subnet is **62**

Activity 4

Subtask 1

```
sudo tcpdump -i any arp -n # flag i means specific interface, flag n means to
disable DNS resolving of hosts
sudo ip neigh flush all # clear local ARP cache
ping 192.168.1.1 # send an ICMP request to the default gateway (192.168.1.1)
```

```
[rocky@rocky-server-3 ~]$ sudo tcpdump -i any arp -n
tcpdump: data link type LINUX_SLL2
dropped privs to tcpdump
tcpdump: verbose output suppressed, use -v[v]... for full protocol decode
listening on any, link-type LINUX_SLL2 (Linux cooked v2), snapshot length 262144 bytes
16:28:40.529878 ens160 B ARP, Request who-has 192.168.1.19 tell 192.168.1.16, length 46
16:28:43.287067 ens160 B ARP, Request who-has 192.168.1.1 tell 192.168.1.16, length 46
16:29:55.951906 ens160 B ARP, Request who-has 192.168.1.32 tell 192.168.1.32, length 46
16:29:56.070333 ens160 B ARP, Request who-has 192.168.1.32 tell 192.168.1.32, length 46
16:30:01.178050 ens160 Out ARP, Request who-has 192.168.1.19 tell 192.168.1.42, length 28
16:30:01.178473 ens160 In ARP, Reply 192.168.1.19 is-at 70:85:c2:68:14:e8, length 46
16:30:56.188048 ens160 B ARP, Request who-has 192.168.1.32 tell 192.168.1.32, length 46
16:31:01.594026 ens160 Out ARP, Request who-has 192.168.1.19 tell 192.168.1.42, length 28
16:31:01.594441 ens160 In ARP, Reply 192.168.1.19 is-at 70:85:c2:68:14:e8, length 46
16:31:02.156086 ens160 B ARP, Request who-has 192.168.1.50 tell 0.0.0.0, length 46
^C
10 packets captured
11 packets received by filter
0 packets dropped by kernel
[rocky@rocky-server-3 ~]$ _

ubuntu@ubuntu-server-1:~$ sudo ip neigh flush all
ubuntu@ubuntu-server-1:~$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1) 56(84) bytes of data.
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.43 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=64 time=1.61 ms
64 bytes from 192.168.1.1: icmp_seq=3 ttl=64 time=1.30 ms
64 bytes from 192.168.1.1: icmp_seq=4 ttl=64 time=1.56 ms
64 bytes from 192.168.1.1: icmp_seq=5 ttl=64 time=1.35 ms
64 bytes from 192.168.1.1: icmp_seq=6 ttl=64 time=1.30 ms
64 bytes from 192.168.1.1: icmp_seq=7 ttl=64 time=1.51 ms
64 bytes from 192.168.1.1: icmp_seq=8 ttl=64 time=1.44 ms
64 bytes from 192.168.1.1: icmp_seq=9 ttl=64 time=1.54 ms
64 bytes from 192.168.1.1: icmp_seq=10 ttl=64 time=1.57 ms
64 bytes from 192.168.1.1: icmp_seq=11 ttl=64 time=1.45 ms
64 bytes from 192.168.1.1: icmp_seq=12 ttl=64 time=1.58 ms
64 bytes from 192.168.1.1: icmp_seq=13 ttl=64 time=1.56 ms
64 bytes from 192.168.1.1: icmp_seq=14 ttl=64 time=1.50 ms
64 bytes from 192.168.1.1: icmp_seq=15 ttl=64 time=1.62 ms
64 bytes from 192.168.1.1: icmp_seq=16 ttl=64 time=1.32 ms
64 bytes from 192.168.1.1: icmp_seq=17 ttl=64 time=1.77 ms
64 bytes from 192.168.1.1: icmp_seq=18 ttl=64 time=1.47 ms
64 bytes from 192.168.1.1: icmp_seq=19 ttl=64 time=1.95 ms
64 bytes from 192.168.1.1: icmp_seq=20 ttl=64 time=1.50 ms
64 bytes from 192.168.1.1: icmp_seq=21 ttl=64 time=1.64 ms
64 bytes from 192.168.1.1: icmp_seq=22 ttl=64 time=1.62 ms
64 bytes from 192.168.1.1: icmp_seq=23 ttl=64 time=1.48 ms
64 bytes from 192.168.1.1: icmp_seq=24 ttl=64 time=1.50 ms
64 bytes from 192.168.1.1: icmp_seq=25 ttl=64 time=1.58 ms
64 bytes from 192.168.1.1: icmp_seq=26 ttl=64 time=1.79 ms
64 bytes from 192.168.1.1: icmp_seq=27 ttl=64 time=1.58 ms
64 bytes from 192.168.1.1: icmp_seq=28 ttl=64 time=1.54 ms
64 bytes from 192.168.1.1: icmp_seq=29 ttl=64 time=1.66 ms
64 bytes from 192.168.1.1: icmp_seq=30 ttl=64 time=1.52 ms
64 bytes from 192.168.1.1: icmp_seq=31 ttl=64 time=1.48 ms
64 bytes from 192.168.1.1: icmp_seq=32 ttl=64 time=1.54 ms
64 bytes from 192.168.1.1: icmp_seq=33 ttl=64 time=1.58 ms
64 bytes from 192.168.1.1: icmp_seq=34 ttl=64 time=1.59 ms
64 bytes from 192.168.1.1: icmp_seq=35 ttl=64 time=1.59 ms
64 bytes from 192.168.1.1: icmp_seq=36 ttl=64 time=1.46 ms
64 bytes from 192.168.1.1: icmp_seq=37 ttl=64 time=1.04 ms
64 bytes from 192.168.1.1: icmp_seq=38 ttl=64 time=2.01 ms
64 bytes from 192.168.1.1: icmp_seq=39 ttl=64 time=1.17 ms
64 bytes from 192.168.1.1: icmp_seq=40 ttl=64 time=1.53 ms
64 bytes from 192.168.1.1: icmp_seq=41 ttl=64 time=1.60 ms
64 bytes from 192.168.1.1: icmp_seq=42 ttl=64 time=1.54 ms
64 bytes from 192.168.1.1: icmp_seq=43 ttl=64 time=1.22 ms
64 bytes from 192.168.1.1: icmp_seq=44 ttl=64 time=1.52 ms
^C
--- 192.168.1.1 ping statistics ---
44 packets transmitted, 44 received, 0% packet loss, time 43080ms
rtt min/avg/max/mdev = 1.036/1.567/3.946/0.395 ms
ubuntu@ubuntu-server-1:~$
```

Subtask 2

```
sudo tcpdump -i ens160 icmp -n  
ping 8.8.8.8 -c5 # send 5 ICMP request to 8.8.8.8
```

```
[rocky@rocky-server-3 ~]$ sudo tcpdump -i ens160 icmp -n  
tcpdump: verbose output suppressed, use -v[v]... for full protocol decode  
listening on ens160, link-type EN10MB (Ethernet), snapshot length 262144 bytes  
16:46:32.138409 IP 192.168.1.16 > 8.8.8.8: ICMP echo request, id 19, seq 1, length 64  
16:46:32.163010 IP 8.8.8.8 > 192.168.1.16: ICMP echo reply, id 19, seq 1, length 64  
16:46:33.139917 IP 192.168.1.16 > 8.8.8.8: ICMP echo request, id 19, seq 2, length 64  
16:46:33.154984 IP 8.8.8.8 > 192.168.1.16: ICMP echo reply, id 19, seq 2, length 64  
16:46:34.141015 IP 192.168.1.16 > 8.8.8.8: ICMP echo request, id 19, seq 3, length 64  
16:46:34.156124 IP 8.8.8.8 > 192.168.1.16: ICMP echo reply, id 19, seq 3, length 64  
16:46:35.143241 IP 192.168.1.16 > 8.8.8.8: ICMP echo request, id 19, seq 4, length 64  
16:46:35.158232 IP 8.8.8.8 > 192.168.1.16: ICMP echo reply, id 19, seq 4, length 64  
16:46:36.145218 IP 192.168.1.16 > 8.8.8.8: ICMP echo request, id 19, seq 5, length 64  
16:46:36.160209 IP 8.8.8.8 > 192.168.1.16: ICMP echo reply, id 19, seq 5, length 64
```

```
ubuntu@ubuntu-server-1:~$ ping 8.8.8.8 -c5  
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data:  
64 bytes from 8.8.8.8: icmp_seq=1 ttl=119 time=24.9 ms  
64 bytes from 8.8.8.8: icmp_seq=2 ttl=119 time=15.8 ms  
64 bytes from 8.8.8.8: icmp_seq=3 ttl=119 time=15.6 ms  
64 bytes from 8.8.8.8: icmp_seq=4 ttl=119 time=15.6 ms  
64 bytes from 8.8.8.8: icmp_seq=5 ttl=119 time=15.5 ms  
  
--- 8.8.8.8 ping statistics ---  
5 packets transmitted, 5 received, 0% packet loss, time 4007ms  
rtt min/avg/max/mdev = 15.491/17.490/24.912/3.712 ms  
ubuntu@ubuntu-server-1:~$
```

Subtask 3

```
sudo tcpdump -i ens160 host neverssl.com and tcp port 80 -A  
# flag A display captured packets in ASCII  
  
# curl tool was used to send HTTP request to neverssl.com  
curl http://neverssl.com
```

No. 7 / 8

```
# ubuntu@ubuntu:~$ cat index.html
<html>
  <head>
    <title>NeverSSL - Connecting ... </title>
    <style>
      body {
        font-family: Montserrat, helvetica, arial, sans-serif;
        font-size: 16px;
        color: #444444;
        margin: 0;
      }
      h2 {
        font-weight: 700;
        font-size: 1.6em;
        margin-top: 30px;
      }
      p {
        line-height: 1.6em;
      }
      .container {
        max-width: 650px;
        margin: 20px auto 20px auto;
        padding-left: 15px;
        padding-right: 15px;
      }
      .header {
        background-color: #42C8FD;
        color: #FFFFFF;
        padding: 10px 0 10px 0;
        font-size: 2.2em;
      }
      .notice {
        background-color: red;
        color: white;
        padding: 10px 0 10px 0;
        font-size: 1.25em;
        animation: flash 4s infinite;
      }
      @keyframes flash {
        0% {
          background-color: red;
        }
        50% {
          background-color: #A80000;
        }
        100% {
          background-color: red;
        }
      }
      /* CSS from Mark Webster https://gist.github.com/markwebster/9bdf30655cdd5279bad13993ac87c85d */
    </style>

    <script>
      var adjectives = [ 'cool', 'calm', 'relaxed', 'soothing', 'serene', 'slow',
        'beautiful', 'wonderful', 'wonderous', 'fun', 'good',
        'glowing', 'inner', 'grand', 'majestic', 'astounding',
        'fine', 'splendid', 'transcendent', 'sublime', 'whole',
        'unique', 'old', 'young', 'fresh', 'clean', 'shiny',
        'shining', 'lush', 'quiet', 'bright', 'silver' ];

      var nouns = [ 'day', 'dawn', 'peace', 'smile', 'love', 'zen', 'laugh',
        'yawn', 'poem', 'song', 'joke', 'verse', 'kiss', 'sunrise',
        'sunset', 'eclipse', 'moon', 'rainbow', 'rain', 'plan',
        'play', 'chart', 'birds', 'stars', 'pathway', 'secret',
        'treasure', 'melody', 'magic', 'spell', 'light', 'morning' ];

      var prefix =
        // Choose 3 zen adjectives
        adjectives.sort(function(){return 0.5-Math.random();}).slice(-3).join('')
        +
        // Coupled with a zen noun
        nouns.sort(function(){return 0.5-Math.random();}).slice(-1).join('');

      window.location.href = 'http://' + prefix + '.neverssl.com/online';
    </script>
  </head>
  <body>
    <noscript>
      <div class="notice">
        <div class="container">
          <div>
            ▲ JavaScript appears to be disabled. NeverSSL's cache-busting works better if you enable JavaScript for <code>neverssl.com/code</code>.
          </div>
        </div>
      </noscript>
      <div class="header">
        <div class="container">
          <h1>NeverSSL</h1>
        </div>
      </div>
      <div class="content">
        <div class="container">
          <h1 id="status"></h1>
          <script>document.querySelector("#status").textContent = "Connecting ...";</script>
          <noscript>
            <h2>What?</h2>
            <p>This website is for when you try to open Facebook, Google, Amazon, etc on a wifi network, and nothing happens. Type "http://neverssl.com" into your browser's url bar, and you'll be able to log on.</p>
            <h2>How?</h2>
            <p>neverssl.com will never use SSL (also known as TLS). No encryption, no strong authentication, no <a href="https://en.wikipedia.org/wiki/HTTP_Strict_Transport_Security">HSTS</a>, no HTTP/2.0, just plain old unencrypted HTTP and forever stuck in the dark ages of internet security.</p>
            <h2>Why?</h2>
            <p>Normally, that's a bad idea. You should always use SSL and secure encryption when possible. In fact, it's such a bad idea that most websites are now using https by default.</p>
            <p>And that's great, but it also means that if you're relying on poorly-behaved wifi networks, it can be hard to get online. Secure browsers and websites using https make it impossible for those wifi networks to send you to a login or payment page. Basically, those networks can't tap into your connection just like attackers can't. Modern browsers are so good that they can remember when a website supports encryption and even if you type in the website name, they'll use https.</p>
            <p>And if the network never redirects you to this page, well as you can see, you're not missing much.</p>
            <a href="https://twitter.com/neverssl">Follow @neverssl</a>
          </noscript>
        </div>
      </div>
    </body>
  </html>
ubuntu@ubuntu:~$
```

Higher resolution photos can be found in the "images" folder.